Report 1: Rudy - A rudimentary web server

Bernardo González Riede

September 14, 2017

1 Introduction

As mentioned in the assignment paper, the purpose of this assignment is to learn:

- *the procedures for using a socket API*

- *the structure of a server process*

- *the HTTP protocol*

The knowledge which is acquired through these concepts mark the fundaments of communication through a network.
Sockets are the base for all communications towards external processes. Sometimes they are even used for communicaion between processes residing on the same machine.
HTTP is a widely used protocol for web servers.

2 Main problems and solutions

One problem encountered was how to simulate different scenarios and participants. Since the exercise is only a small web server, all test could have been realized on the same machine (running Windows 10). Still, the thirst for performing real world test with physical networks and GNU/Linux servers was far too attractive. Therefore the test involved three participants:

Name	CPU	RAM	Location	OS
clientLocal	4-core	4GB	KTH Kista	Windows 10
server	1-core	512MB	Frankfurt	Ubuntu-server 16.04
clientRemote	1-core	512MB	Frankfurt	Ubuntu-server 16.04

Table 1: Information about the participants

The two Ubuntu machines were hosted on Digital Ocean's Datacenter in Frankfurt, Germany. It's a virtualized datacenter, hence the incomplete description about their hardware.

3 Evaluation

5 scenarios were used to test the web server.

- A: Rudy & test executed on clientLocal.

- B: Rudy & test executed on server.

- C: Rudy on server & test on clientLocal.

- D: Rudy on server & test on clientRemote.

- E: Rudy on server while test was executed on both clients simulatenously.

The following table results from doing 100 request to the server. An artificial delay of 40ms was introduced in a second run to simulate file handling.


Figure 1: Comparision of scenarios and artificial delay introduced.

This shows a clear impact of the artificial delay, since it's adds a constant 4s to the 100 requests. Scenario E is special and has to be interpreted. In the case without an artificial delay, it's impact is $<1\%$, but sometimes more noticeable. Contrary to this, running the server with an artificial delay between request elevates the time it takes to complete by up to 50%.

3.1 Multiprocessor tests

Several test were made on clientLocal (4 logical cores) using the files cre-
ated in the assignment without modifications and the provided *rudy4.erl &
test2.erl*. Several combinations running an erlang shell with 1,2 & 4 cores
support were done. This is achieved by executing the shell with *-smp en-
abled -S N* where *N* is the desired amout of simultaneous processes being
supported.
Nonetheless, the observations resultet in no significant change. While the
artificial delay, running 4x 10 request, accounts for 1600ms, the complete
benchmark always endend between 1985ms and 2016ms. The benchmark
was executed in different erlang shells, too.

4 Conclusions

The problem gave some good ideas about how to implement a server to
respond via HTTP instead of the normal messagin of Erlang. This gives
a more universal approach since sockets and HTTP are in many, if not all
progamming languages, present.